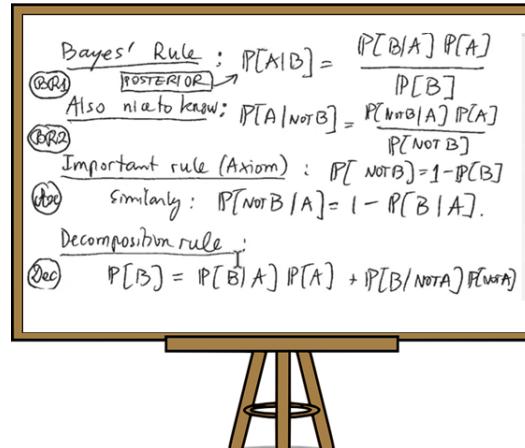


Bridging Theory and Practice Part 2: Intro to Bayesian computation

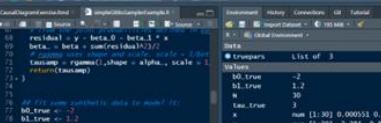
Katherine Muller



Theory



Bayes Rule!



The screenshot shows a Shiny application running in a browser window titled "shinyapp". The application has a sidebar with "Data" and "Model" tabs, and a main panel with "Residuals" and "Fits" sections. The code in the editor is as follows:

```
residual <- y - beta_0 - beta_1 * x
beta_0 <- mean(y)
beta_1 <- cov(x, residual) / var(x)
alpha <- rnorm(1, mean(residual), scale = 1)
return(c(beta_0, beta_1, alpha))
}

# If you run this function with synthetic data in mode == "fit"
# bl_true <- 2
# bl1_true <- 1.2
# tau_true <- 0.5

trueparams <- list(bl_true = bl_true,
                      beta_1 = bl1_true,
                      tau_true = tau_true)

ui <- fluidPage(
  titlePanel("Simple Linear Regression"),
  sidebarLayout(
    sidebarPanel("Data", "Model"),
    mainPanel("Residuals", "Fits")
  )
)
```

Computation

Recap of Part 1

- Statistical inference involves estimating **unknown** or **unobservable parameters** from **observable data** (e.g., **disease resistance** from **disease occurrence**).
- Frequentist vs. Bayesian inference.
- Simple linear regression model for rainfed rice yield vs. rainfall.
- Without computer sampling, Bayesian inference can only be performed on **simple models** with **very restricted priors**.

Goals for Part 2

- Build your intuitive understanding about how Bayesian theory extends to Bayesian computation
- Gain familiarity with the basics of Bayesian computation using simple linear regression model

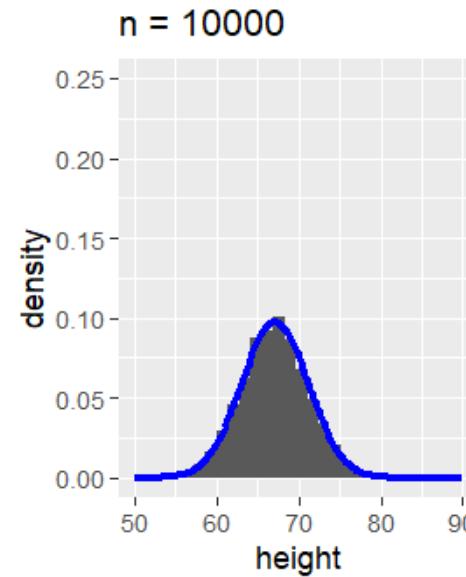
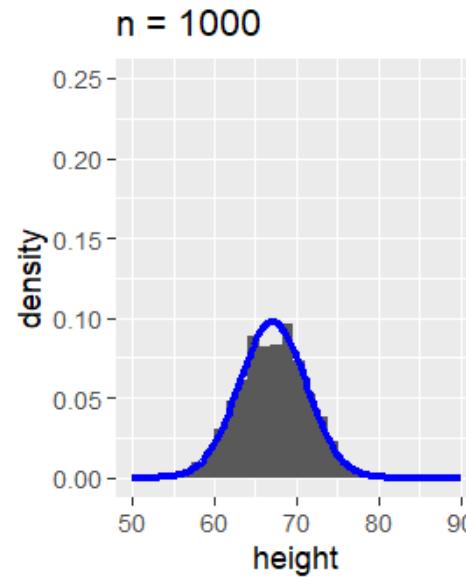
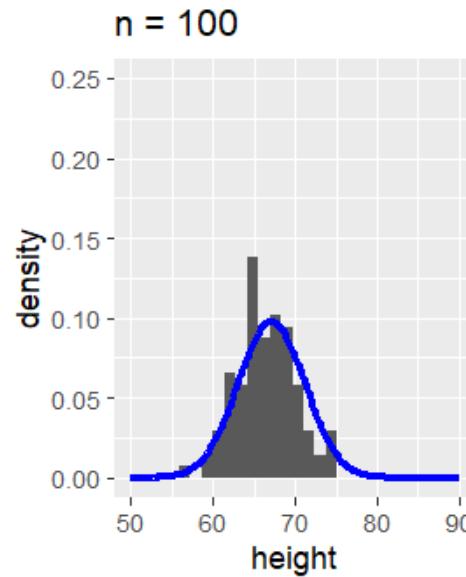
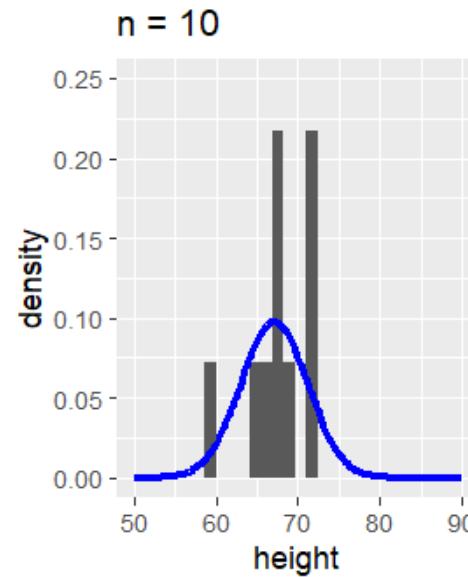
Computers turn an analytical task into a sampling task

$\text{height} \sim \mathcal{N}(\mu, \sigma)$

a.k.a. Law of Large Numbers

`rnorm(n, mean, sd)`

More samples → closer to theoretical distribution



Markov chain Monte Carlo (MCMC)

- Family of computer algorithms used for Bayesian data analysis
- Approximate theoretical distributions through iterative sampling
- Gibbs sampling is one type of MCMC algorithm

Building our model



1. Specify the likelihood
2. Specify priors
3. Sample posterior

$$P(\boldsymbol{\theta} | \mathcal{D}) = \frac{P(\mathcal{D} | \boldsymbol{\theta}) P(\boldsymbol{\theta})}{P(\mathcal{D})}$$

Building our model



1. Specify the likelihood

Choose a **likelihood function** that represents our data-generating process

2. Specify priors

$$y_i | x_i, \beta_0, \beta_1, \tau \sim \mathcal{N}(\beta_0 + \beta_1 x_i, 1/\tau)$$

3. Sample posterior

Yield data (y) are generated on a normal distribution with the mean influenced by rainfall (x).

$$P(\theta|D) = \frac{P(D|\theta)P(\theta)}{P(D)}$$

We can simulate y by sampling a normal distribution .

```
b0 <- 0; b1 <- 1.2; tau <- 3 # "true" parameters  
N <- 30 # number of observations of x and y  
x <- runif(N, -1, 1) ## random number between -1 and 1  
y <- rnorm(N, mean = b0 + b1*x, sd = 1/sqrt(tau))
```

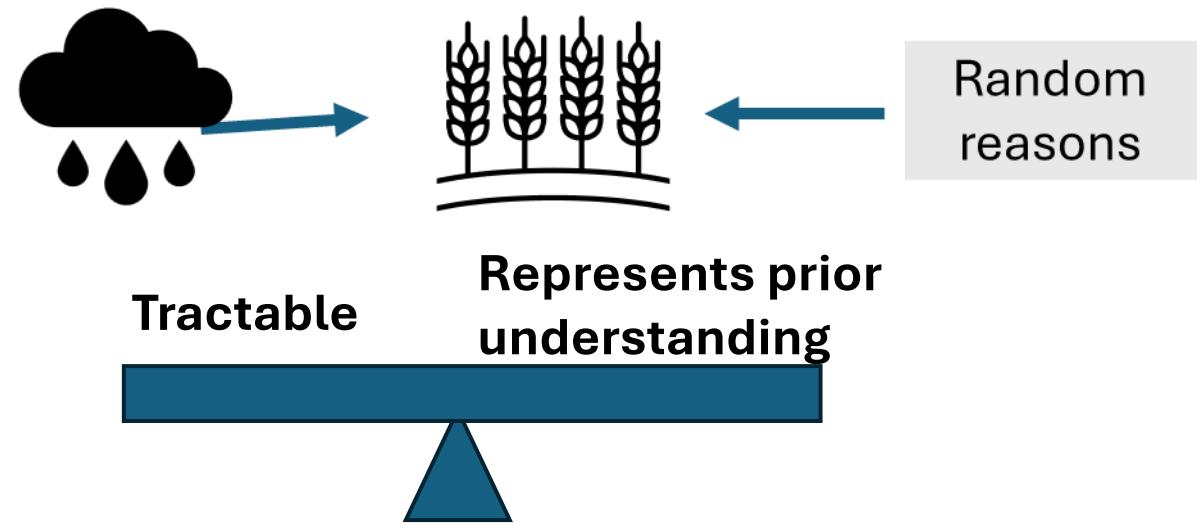
Building our model

1. Specify the likelihood

2. Specify priors

3. Sample posterior

$$P(\theta|D) = \frac{P(D|\theta)P(\theta)}{P(D)}$$



Building our model

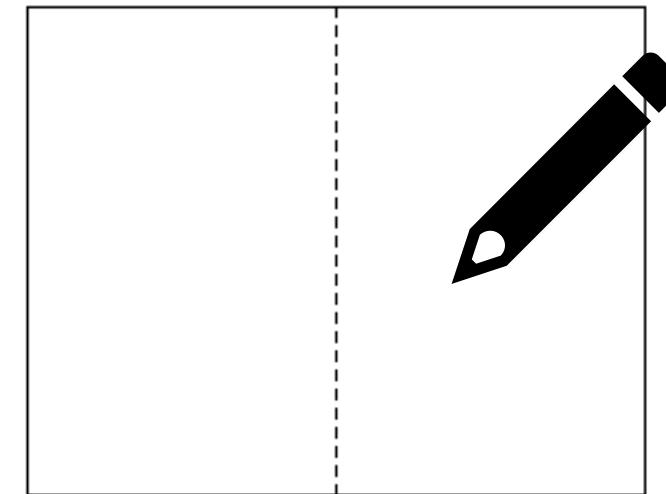
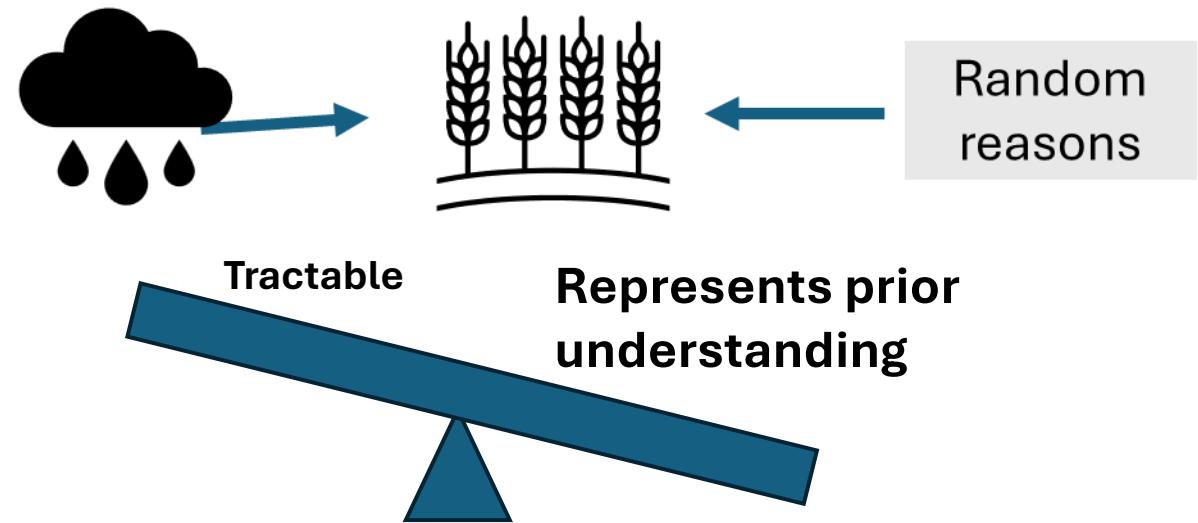
1. Specify the likelihood

2. Specify priors

3. Sample posterior

$$P(\theta|D) = \frac{P(D|\theta)P(\theta)}{P(D)}$$

Plausibility of
parameter
value $P(\theta_j)$



θ_j

parameter value

Building our model

1. Specify the likelihood

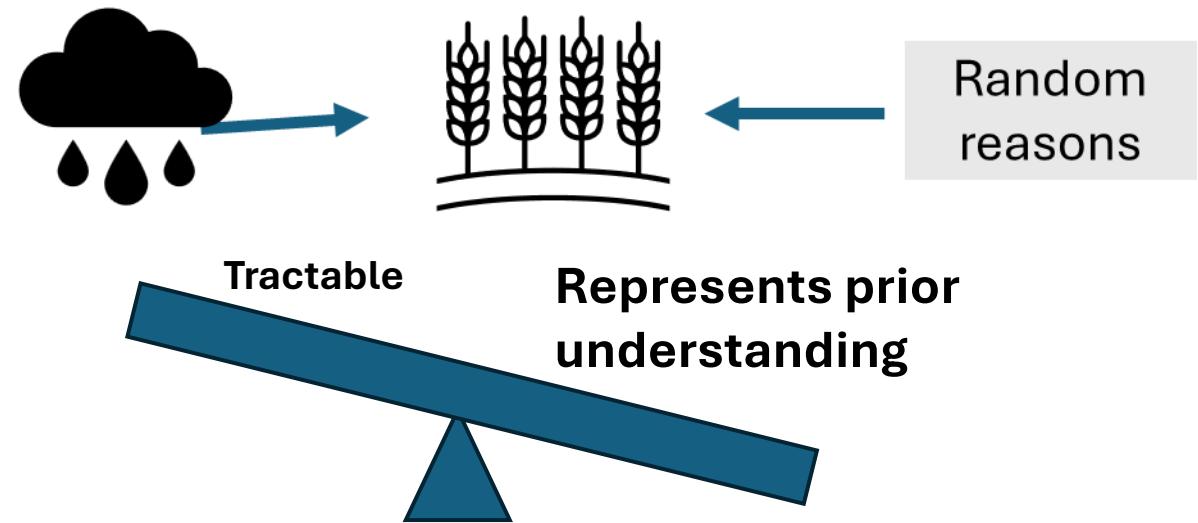
2. Specify priors

3. Sample posterior

$$P(\theta|D) = \frac{P(D|\theta)P(\theta)}{P(D)}$$



I know that I know
nothing!
- Plato



Building our model

1. Specify the likelihood

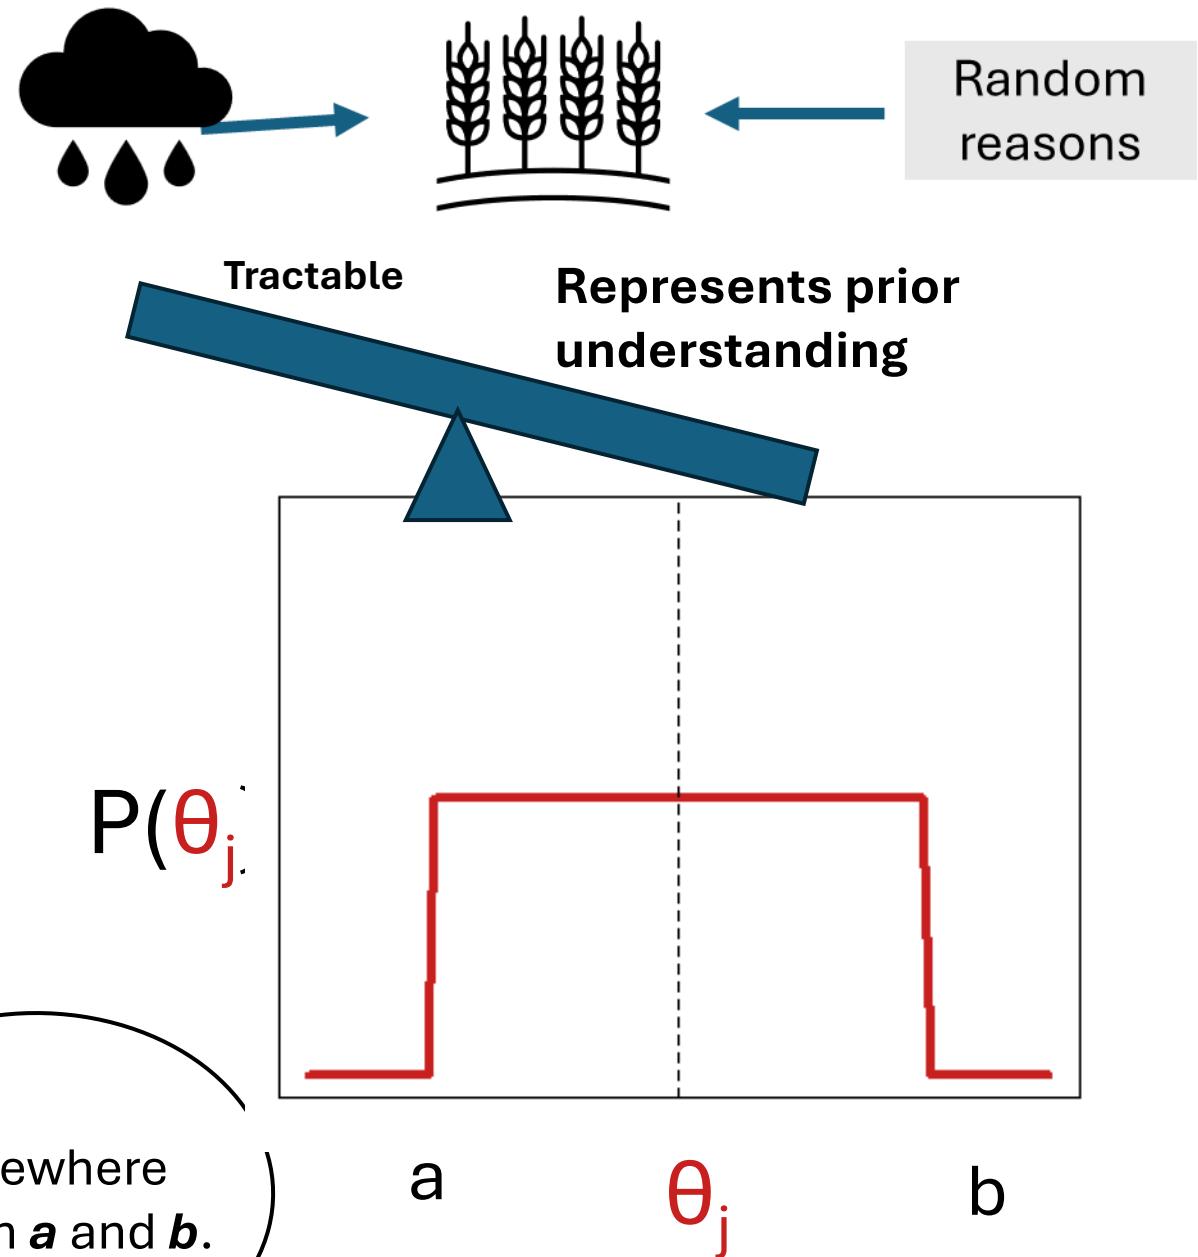
2. Specify priors

3. Sample posterior

$$P(\theta|D) = \frac{P(D|\theta)P(\theta)}{P(D)}$$



It's somewhere
between a and b .



Building our model



1. Specify the likelihood

Data generating function

2. Specify priors

3. Sample posterior

$$P(\boldsymbol{\theta}|D) = \frac{P(D|\boldsymbol{\theta})P(\boldsymbol{\theta})}{P(D)}$$

Tractable

Represents prior understanding

Choose prior functions that will combine with the likelihood function to **create posteriors we can easily sample.**

Parameter generating functions

$$P(\boldsymbol{\theta})$$

$$\theta_j | \theta_j^* \sim f(\theta_j^*)$$

$$P(\boldsymbol{\theta}|D)$$

$$\theta_j | \theta_j^*, D \sim f(\theta_j^*, D)$$

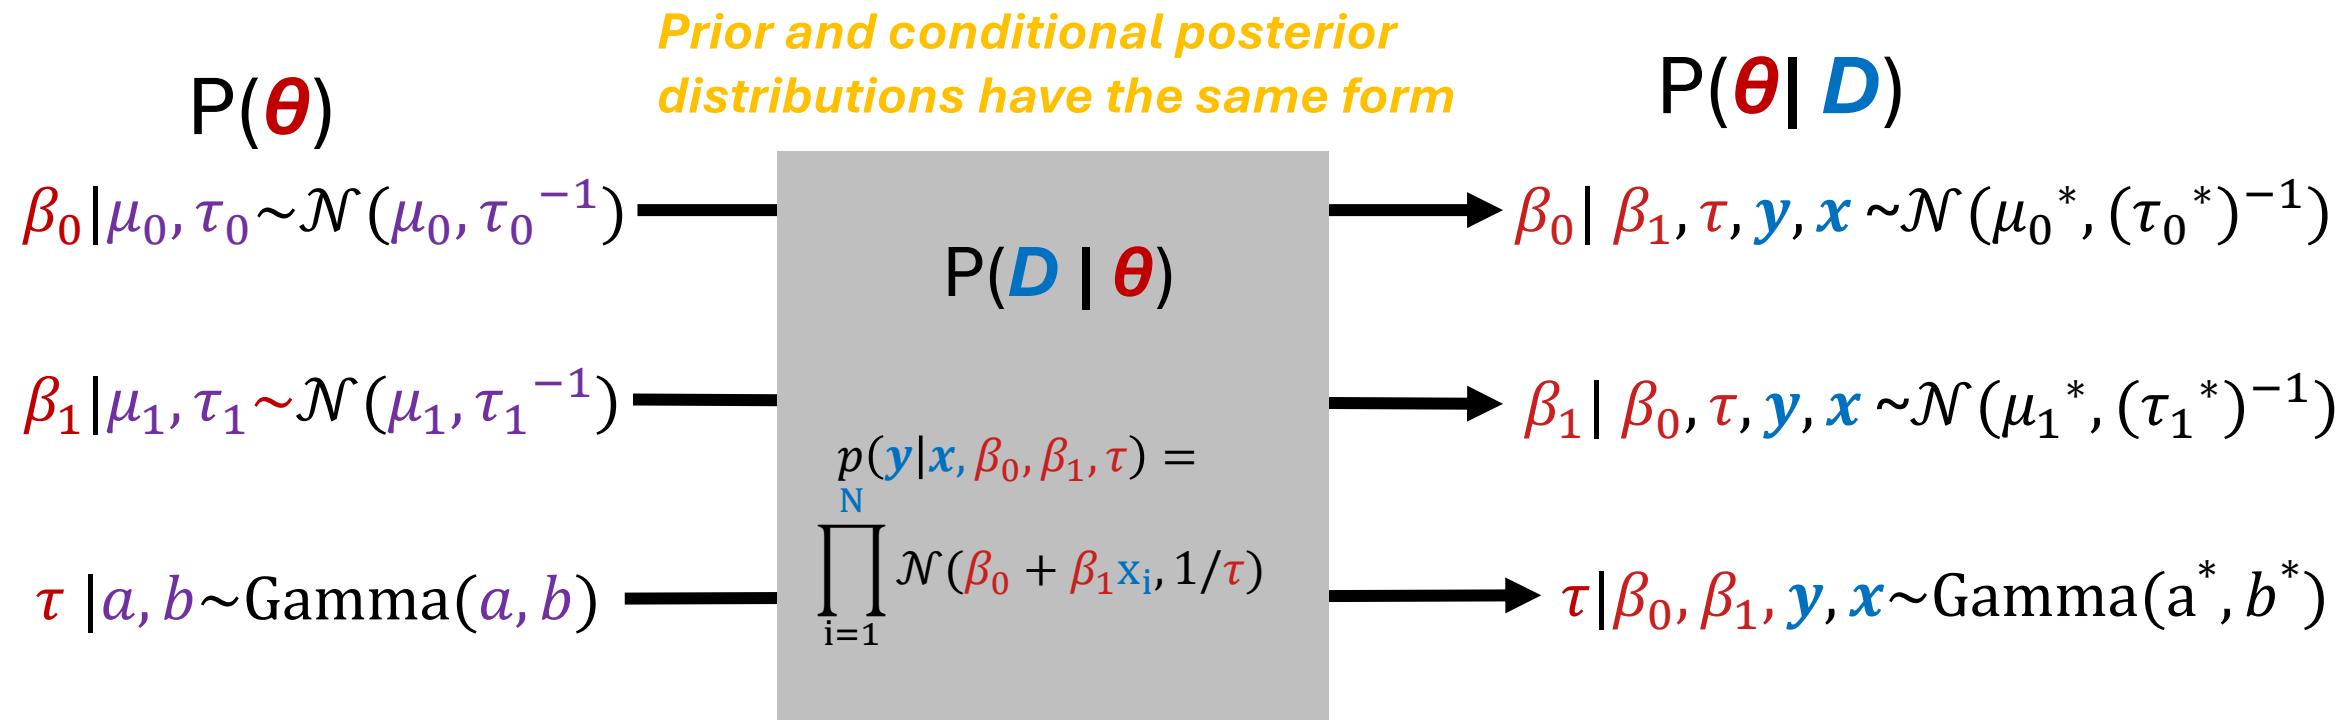
Same type of probability distribution

$$\theta_j \sim \mathcal{N}(\mu_j, \tau_j^{-1})$$

$$\theta_j | D \sim \mathcal{N}(\mu_j^*, \tau_j^{*-1})$$

Conditionally conjugate prior

Conditionally conjugate priors for simple linear regression



- Intuitive for Bayesian updating
- Can speed up processing
- Good default if you don't have strong prior understanding
- May not adequately model for prior understanding in all situations
- Not required for Bayesian computation

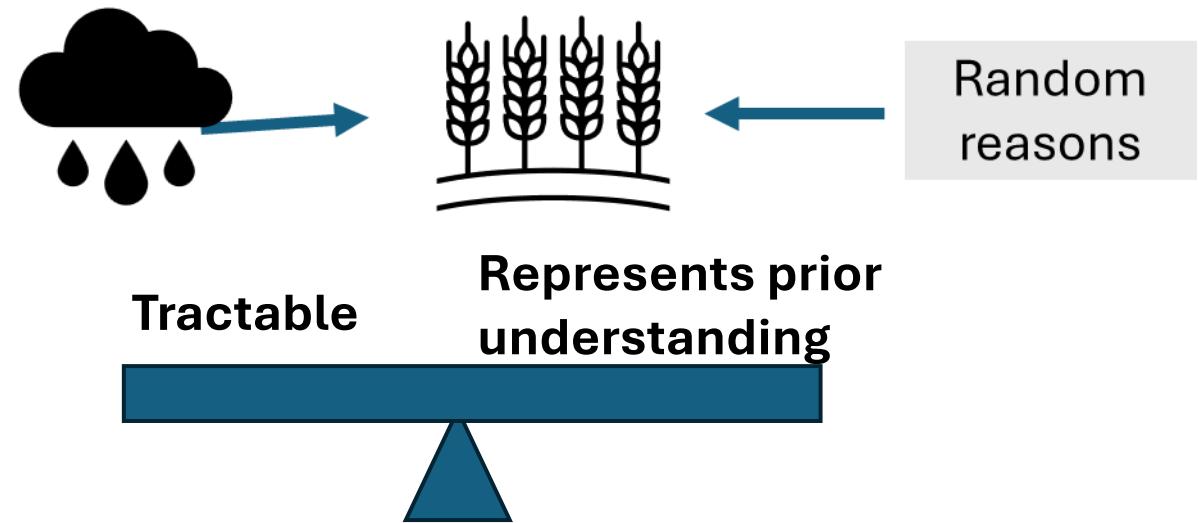
Building our model

1. Specify the likelihood

2. Specify priors

3. Sample posterior

$$P(\theta|D) = \frac{P(D|\theta)P(\theta)}{P(D)}$$

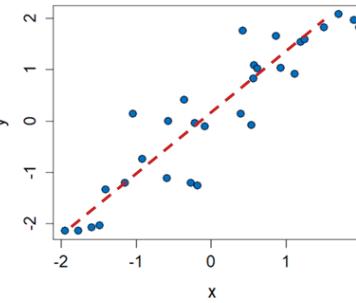


Use **tractable priors** (e.g., conditionally conjugate) and adjust **hyperparameters** to **reflect prior understanding**.

$$P(\theta|D) = \frac{P(D|\theta)P(\theta)}{P(D)}$$

Simple linear regression:

$$\begin{aligned}y &= \beta_0 + \beta_1 x + \varepsilon \\ \varepsilon &\sim \mathcal{N}(0, 1/\tau)\end{aligned}$$



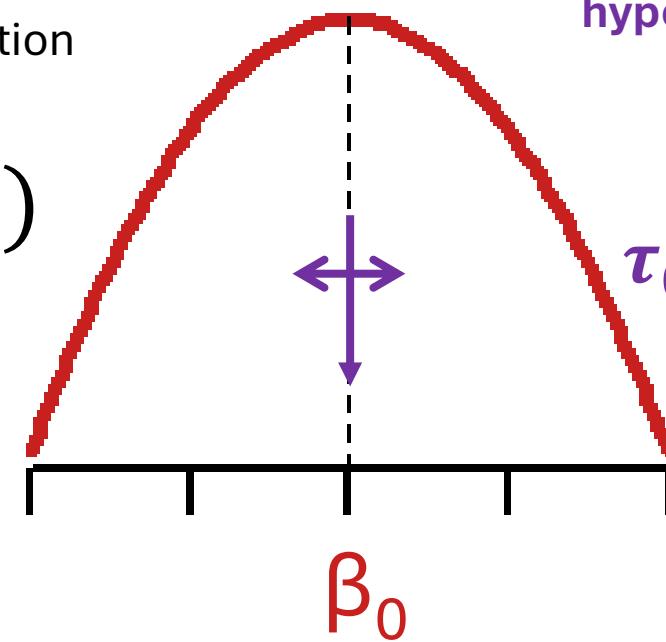
Intercept



Since we standardized x and y , the intercept is probably close to zero

Conjugate \rightarrow Normal distribution

$$P(\beta_0)$$



hyperparameters

$$\mu_0 = 0$$

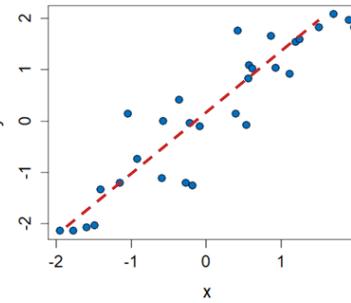
Some number expressing our uncertainty

$$\tau_0^{-1}$$

$$P(\theta|D) = \frac{P(D|\theta)P(\theta)}{P(D)}$$

Simple linear regression:

$$\begin{aligned}y &= \beta_0 + \beta_1 x + \varepsilon \\ \varepsilon &\sim \mathcal{N}(0, 1/\tau)\end{aligned}$$



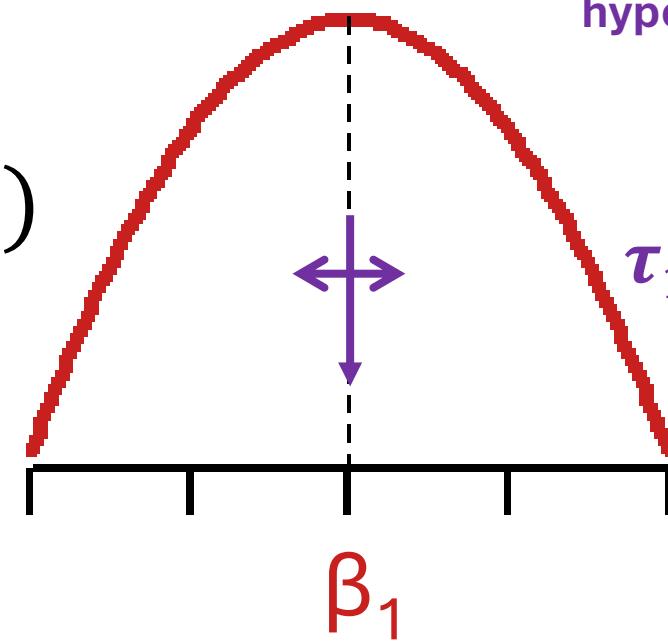
Slope



The slope of x vs. y could be negative, positive, or zero.

Conjugate \rightarrow Normal distribution

$$P(\beta_1)$$



hyperparameters

$$\mu_1 = 0$$

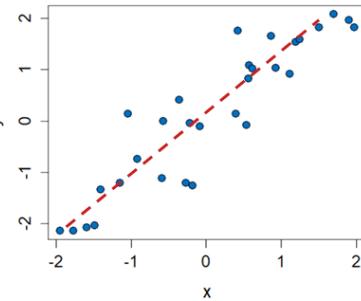
Some number expressing our uncertainty

$$\tau_1^{-1} =$$

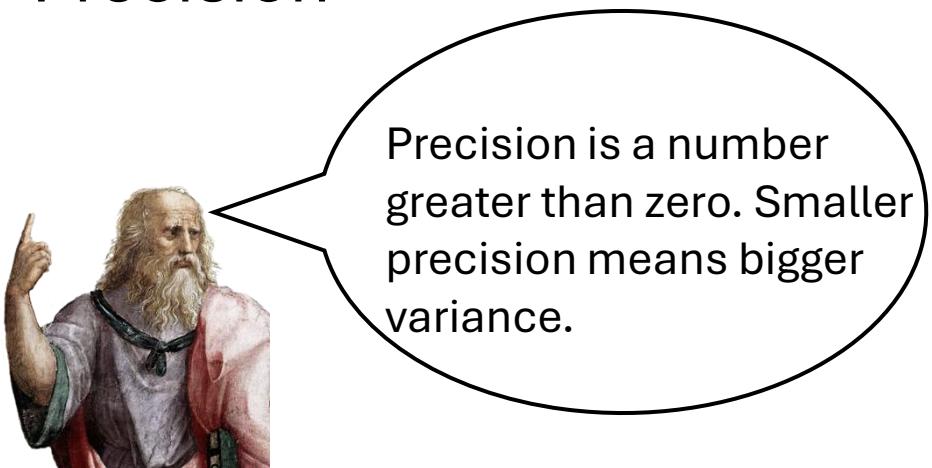
$$P(\theta|D) = \frac{P(D|\theta)P(\theta)}{P(D)}$$

Simple linear regression:

$$\begin{aligned}y &= \beta_0 + \beta_1 x + \varepsilon \\ \varepsilon &\sim \mathcal{N}(0, 1/\tau)\end{aligned}$$



Precision



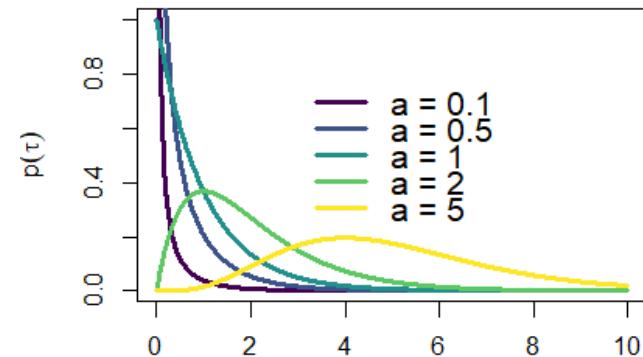
Precision is a number greater than zero. Smaller precision means bigger variance.

Conjugate \rightarrow Gamma distribution

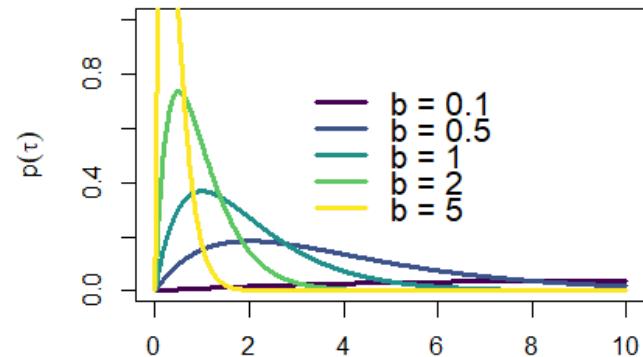
a = shape

b = rate

Shape effect (a), Rate (b) = 1

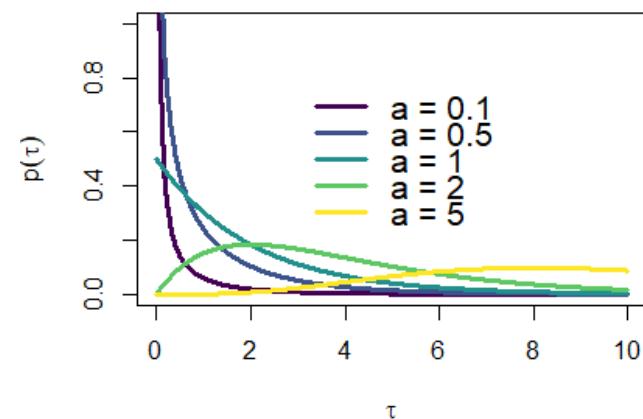


Rate effect (b), Shape (a) = 2

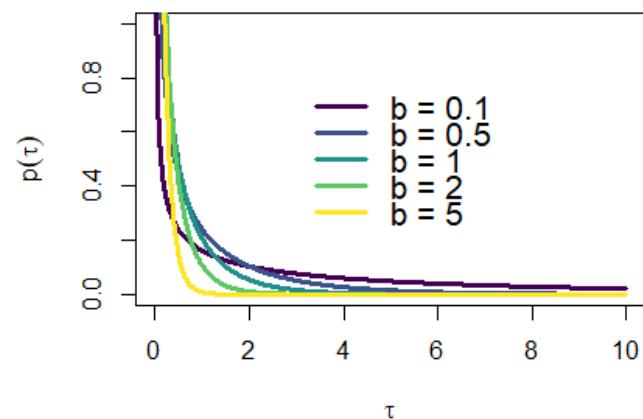


Tricky—hyperparameter choices affect posterior sampling

Shape effect (a), Rate (b) = 0.5



Rate effect (b), Shape (a) = 0.5



Building our model

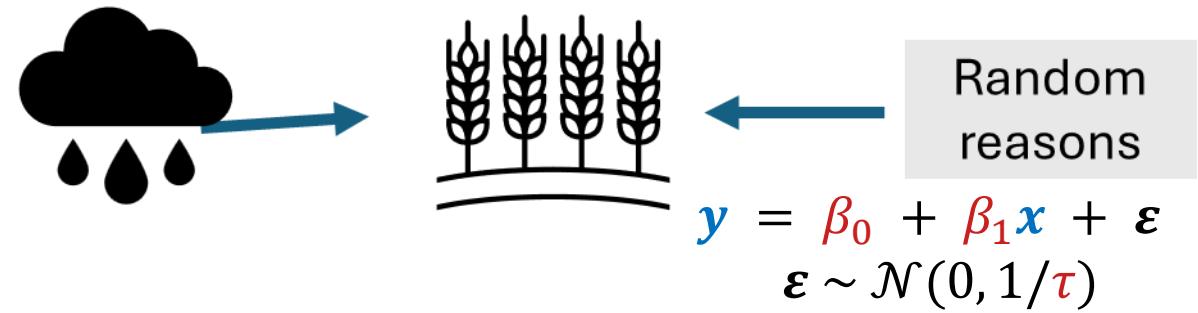
1. Specify the likelihood

2. Specify priors

3. Sample posterior

$$P(\theta|D) = \frac{P(D|\theta)P(\theta)}{P(D)}$$

Normalizing
constant



Sample **one parameter at a time** conditional on the **data** and **other parameters**

Full conditional posterior

Likelihood	Prior
$p(\beta_0 \beta_1, \tau, y, x) \propto p(y, x \beta_0, \beta_1, \tau)$	$p(\beta_0)$
$p(\beta_1 \beta_0, \tau, y, x) \propto p(y, x \beta_0, \beta_1, \tau)$	$p(\beta_1)$
$p(\tau \beta_0, \beta_1, y, x) \propto p(y, x \beta_0, \beta_1, \tau)$	$p(\tau)$

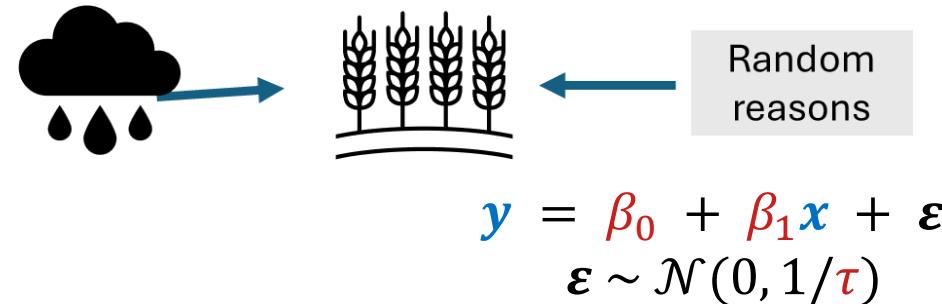
Gibbs sampling!

Gibbs sampling

- A **computer algorithm** used to **approximate the full posterior distribution** of one or more parameters
- Works by **iteratively sampling the full conditional distributions** of each parameter
- One member of the broader family of **Markov Chain Monte Carlo (MCMC)** methods

To build our Gibbs sampler, we need...

1. Data on rainfall and rice yield.



2. A data-generating model

3. A likelihood function representing our data-generating model

$$p(\mathbf{y}, \mathbf{x} | \beta_0, \beta_1, \tau) = \prod_{i=1}^N \mathcal{N}(\beta_0 + \beta_1 x_i, 1/\tau)$$

4. Prior distributions and hyperparameters for each parameter

intercept	$p(\beta_0) = \mathcal{N}(\mu_0, \tau_0^{-1})$
slope	$p(\beta_1) = \mathcal{N}(\mu_1, \tau_1^{-1})$
precision	$p(\tau) = \text{Gamma}(a, b)$

5. Full conditional posterior distributions for each parameter

Calculate using Bayes' Rule!

Hyperparameters	$\mu_0 = 0$	$\tau_0 = 1$
	$\mu_1 = 0$	$\tau_1 = 1$
	$a = 2$	$b = 1$

Full conditional posteriors via Bayes' Rule

Intercept β_0

$$p(\beta_0 | \beta_1, \tau, \mathbf{y}, \mathbf{x}) \propto p(\mathbf{y}, \mathbf{x} | \beta_0, \beta_1, \tau) p(\beta_0)$$

$$= \prod_{i=1}^N \mathcal{N}(\beta_0 + \beta_1 x_i, 1/\tau) \mathcal{N}(\mu_0, \tau_0^{-1})$$

$$= \prod_{i=1}^N \sqrt{\frac{\tau}{2\pi}} e^{-\frac{\tau}{2}(y_i - (\beta_0 + \beta_1 x_i))^2} \sqrt{\frac{\tau_0}{2\pi}} e^{-\frac{\tau_0}{2}(\beta_0 - \mu_0)^2}$$

*Some clever rearranging **

$$\beta_0 | \beta_1, \tau, \mathbf{y}, \mathbf{x} \sim \mathcal{N}\left(\frac{\tau_0 \mu_0 + \tau \sum_{i=1}^N (y_i - \beta_1 x_i)}{\tau_0 + \tau N}, (\tau_0 + \tau N)^{-1}\right)$$

Normal probability density function

$$z | \mu, \tau \sim \mathcal{N}(\mu, \tau^{-1})$$

$$p(z|\mu, \tau) = \mathcal{N}(\mu, \tau^{-1})$$

$$p(z|\mu, \tau) = \sqrt{\frac{\tau}{2\pi}} e^{-\frac{\tau}{2}(z-\mu)^2}$$

* Read this for a full breakdown of the steps <https://stmorse.github.io/journal/gibbs.html>

Full conditional posteriors via Bayes' Rule

Slope β_1

$$\begin{aligned} p(\beta_1 | \beta_0, \tau, \mathbf{y}, \mathbf{x}) &\propto p(\mathbf{y}, \mathbf{x} | \beta_0, \beta_1, \tau) p(\beta_1) \\ &= \prod_{i=1}^N \mathcal{N}(\beta_0 + \beta_1 x_i, 1/\tau) \mathcal{N}(\mu_1, \tau_1^{-1}) \\ &= \left[\prod_{i=1}^N \sqrt{\frac{\tau}{2\pi}} e^{-\frac{\tau}{2}(y_i - (\beta_0 + \beta_1 x_i))^2} \sqrt{\frac{\tau_1}{2\pi}} e^{-\frac{\tau_1}{2}(\beta_1 - \mu_1)^2} \right] \end{aligned}$$

Some clever rearranging *

$$\beta_1 | \beta_0, \tau, \mathbf{y}, \mathbf{x} \sim \mathcal{N}\left(\frac{\tau_1 \mu_1 + \tau \sum_{i=1}^N (y_i - \beta_0) x_i}{\tau_1 + \tau \sum_{i=1}^N x_i^2}, (\tau_1 + \tau \sum_{i=1}^N x_i^2)^{-1}\right)$$

Normal probability density function

$$z | \mu, \tau \sim \mathcal{N}(\mu, \tau^{-1})$$

$$p(z|\mu, \tau) = \mathcal{N}(\mu, \tau^{-1})$$

$$p(z|\mu, \tau) = \sqrt{\frac{\tau}{2\pi}} e^{-\frac{\tau}{2}(z-\mu)^2}$$

* Read this for a full breakdown of the steps <https://stmorse.github.io/journal/gibbs.html>

Full conditional posteriors via Bayes' Rule

Precision τ

$$\begin{aligned} p(\tau | \beta_0, \beta_1, \mathbf{y}, \mathbf{x}) &\propto p(\mathbf{y}, \mathbf{x} | \beta_0, \beta_1, \tau) p(\tau) \\ &= \prod_{i=1}^N \mathcal{N}(\beta_0 + \beta_1 x_i, 1/\tau) \text{Gamma}(a, b) \\ &= \prod_{i=1}^N \sqrt{\frac{\tau}{2\pi}} e^{-\frac{\tau}{2}(y_i - (\beta_0 + \beta_1 x_i))^2} \frac{b^a}{\Gamma(a)} \tau^{a-1} e^{-b\tau} \end{aligned}$$


*Some clever rearranging ** 

$$\tau | \beta_0, \beta_1, \mathbf{y}, \mathbf{x} \sim \text{Gamma}\left(a + \frac{N}{2}, b + \frac{1}{2} \sum_{i=1}^N (y_i - (\beta_0 + \beta_1 x_i))^2\right)$$

Gamma probability density function

$$x | a, b \sim \text{Gamma}(a, b)$$

$$p(x|a, b) = \text{Gamma}(a, b)$$

$$p(x|a, b) = \frac{b^a}{\Gamma(a)} x^{a-1} e^{-bx}$$

* Read this for a full breakdown of the steps <https://stmorse.github.io/journal/gibbs.html>

We can code the full conditional posterior for a single parameter as a sampling function in R!

$$\text{Intercept } \beta_0 | \beta_1, \tau, y, x \sim \mathcal{N}\left(\frac{\tau_0 \mu_0 + \tau \sum_{i=1}^N (y_i - \beta_1 x_i)}{\tau_0 + \tau N}, (\tau_0 + \tau)^{-1}\right)$$

```
sample_b0 <- function(x,y, beta_1, tau, mu_0, tau_0) {
  N <- length(x)
  new_precision <- tau_0 + tau * N
  new_mean <- (tau_0 * mu_0 + tau * sum(y - beta_1 * x)) / new_precision
  b0samp <- rnorm(1,new_mean, 1/sqrt(new_precision))
  return(b0samp)
}
```

Generates one value of β_0 by sampling on a normal distribution with mean and sd calculated from input values for the **data, **parameters**, and **hyperparameters****

We can code the full conditional posterior for a single parameter as a sampling function in R!

Slope $\beta_1 | \beta_0, \tau, \mathbf{y}, \mathbf{x} \sim \mathcal{N}\left(\frac{\tau_1 \mu_1 + \tau \sum_{i=1}^N (y_i - \beta_0) x_i}{\tau_1 + \tau \sum_{i=1}^N x_i^2}, (\tau_1 + \tau \sum_{i=1}^N x_i^2)^{-1}\right)$

```
sample_b1 <- function(x,y, beta_0, tau, mu_1, tau_1){  
  N <- length(x)  
  new_precision <- tau_0 + tau * sum(x^2)  
  new_mean <- (tau_1 * mu_1 + tau * sum(x*(y - beta_0)))/new_precision  
  b1samp <- rnorm(1,new_mean, 1/sqrt(new_precision))  
  return(b1samp)  
}
```

Generates one value of β_1 by **sampling on a normal distribution with mean and sd **calculated from input values for the data, parameters, and hyperparameters****

We can code the full conditional posterior for a single parameter as a sampling function in R!

Precision $\tau | \beta_0, \beta_1, y, x \sim \text{Gamma}\left(a + \frac{N}{2}, b + \frac{1}{2} \sum_{i=1}^N (y_i - (\beta_0 + \beta_1 x_i))^2\right)$

```
sample_tau <- function(x,y, beta_0, beta_1, a, b){  
  N <- length(x)  
  new_shape <- a + N/2  
  residual <- y - beta_0 - beta_1 * x  
  new_rate <- b + sum(residual^2)/2  
  tausamp <- rgamma(1, shape = new_shape, rate = new_rate)  
  return(tausamp)  
}
```

Generates one value of τ by *sampling on a normal distribution* with shape and rate calculated from input values for the **data, **parameters**, and **hyperparameters****

Gibbs sampling algorithm

1. Set a starting value
for each parameter

2. Sample each
parameters using input
data and previous values
of other parameters

3. Repeat for many
iterations

Initialization (iteration 0):

$$\beta_0^{(0)} = 0, \beta_1^{(0)} = 0, \tau^{(0)} = 1$$

Iteration 1:

Update intercept: $\beta_0^{(1)} = \text{sample_b0}(\beta_1^{(0)}, \tau^{(0)}, \mathbf{x}, \mathbf{y})$

Update slope: $\beta_1^{(1)} = \text{sample_b1}(\beta_0^{(1)}, \tau^{(0)}, \mathbf{x}, \mathbf{y})$

Update precision: $\tau^{(1)} = \text{sample_tau}(\beta_0^{(1)}, \beta_1^{(1)}, \mathbf{x}, \mathbf{y})$

Iteration t:

Update intercept: $\beta_0^{(t)} = \text{sample_b0}(\beta_1^{(t-1)}, \tau^{(t-1)}, \mathbf{x}, \mathbf{y})$

Update slope: $\beta_1^{(t)} = \text{sample_b1}(\beta_0^{(t)}, \tau^{(t-1)}, \mathbf{x}, \mathbf{y})$

Update precision: $\tau^{(t)} = \text{sample_tau}(\beta_0^{(t)}, \beta_1^{(t)}, \mathbf{x}, \mathbf{y})$

Next: Look through the Gibbs sampler code and exercises in R.

Runnable code:
GibbsSampler.Rmd

Formatted output:
GibbsSampler.html

